

```
In [2]: import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix
```

```
In [3]: df = pd.read_csv('/content/drive/MyDrive/survey_lung_cancer.csv')
```

```
In [4]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 309 entries, 0 to 308
Data columns (total 16 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   GENDER                                309 non-null    object
1   AGE                                   309 non-null    int64
2   SMOKING                              309 non-null    int64
3   YELLOW_FINGERS                       309 non-null    int64
4   ANXIETY                              309 non-null    int64
5   PEER_PRESSURE                        309 non-null    int64
6   CHRONIC_DISEASE                      309 non-null    int64
7   FATIGUE                             309 non-null    int64
8   ALLERGY                              309 non-null    int64
9   WHEEZING                             309 non-null    int64
10  ALCOHOL_CONSUMING                    309 non-null    int64
11  COUGHING                             309 non-null    int64
12  SHORTNESS_OF_BREATH                  309 non-null    int64
13  SWALLOWING_DIFFICULTY                309 non-null    int64
14  CHEST_PAIN                           309 non-null    int64
15  LUNG_CANCER                          309 non-null    object
dtypes: int64(14), object(2)
memory usage: 38.8+ KB
```

```
In [5]: df.head()
```

```
Out[5]:
```

	GENDER	AGE	SMOKING	YELLOW_FINGERS	ANXIETY	PEER_PRESSURE	CH DI
0	M	69	1	2	2	1	
1	M	74	2	1	1	1	
2	F	59	1	1	1	2	
3	M	63	2	2	2	1	
4	F	63	1	2	1	1	

```
In [6]: # Perform one-hot encoding on the 'gender' column
# Encode the 'GENDER' column: Male ('M') as 1 and Female ('F') as 0
df['GENDER'] = df['GENDER'].apply(lambda x: 1 if x == 'M' else 0)
```

```
In [7]: df.head()
```

Out[7]:

	GENDER	AGE	SMOKING	YELLOW_FINGERS	ANXIETY	PEER_PRESSURE	CH DI
<b>0</b>	1	69	1	2	2	1	
<b>1</b>	1	74	2	1	1	1	
<b>2</b>	0	59	1	1	1	2	
<b>3</b>	1	63	2	2	2	1	
<b>4</b>	0	63	1	2	1	1	

```
In [8]: # check for duplicate records
df.duplicated().sum()
```

Out[8]: 33

```
In [9]: # Drop duplicates records
df.drop_duplicates(inplace=True)
```

```
In [10]: df.duplicated().sum()
```

Out[10]: 0

```
In [12]: # check for null values
df.isna().sum()
```

Out[12]:

	0
<b>GENDER</b>	0
<b>AGE</b>	0
<b>SMOKING</b>	0
<b>YELLOW_FINGERS</b>	0
<b>ANXIETY</b>	0
<b>PEER_PRESSURE</b>	0
<b>CHRONIC DISEASE</b>	0
<b>FATIGUE</b>	0
<b>ALLERGY</b>	0
<b>WHEEZING</b>	0
<b>ALCOHOL CONSUMING</b>	0
<b>COUGHING</b>	0
<b>SHORTNESS OF BREATH</b>	0
<b>SWALLOWING DIFFICULTY</b>	0
<b>CHEST PAIN</b>	0
<b>LUNG_CANCER</b>	0

**dtype:** int64

Logistic Regression Model

```
In [13]: # Initialize the logistic regression model
logistic_model = LogisticRegression(max_iter=1000, random_state=42)
```

```
In [14]: # Feature Selection and Engineering
# Identifying features (X) and target variable (y)
X = df.drop(columns=['LUNG_CANCER'], axis=1)
y = df['LUNG_CANCER']
```

```
In [15]: # Train-Test Split
# Splitting the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, ran
```

```
In [16]: # Train the model on the training data
logistic_model.fit(X_train, y_train)
```

Out[16]:

▼

LogisticRegression

LogisticRegression(max\_iter=1000, random\_state=42)

```
In [17]: # Make predictions on the test set
y_pred = logistic_model.predict(X_test)
```

```
In [18]: # Calculate the accuracy
accuracy = accuracy_score(y_test, y_pred)
print(accuracy)
```

0.891566265060241

```
In [19]: # Confusion matrix
confusion_matrix(y_test, y_pred)
```

```
Out[19]: array([[ 4,  9],
                [ 0, 70]])
```

```
In [20]: # Check accuracy of model
print(classification_report(y_test,y_pred))
```

	precision	recall	f1-score	support
NO	1.00	0.31	0.47	13
YES	0.89	1.00	0.94	70
accuracy			0.89	83
macro avg	0.94	0.65	0.71	83
weighted avg	0.90	0.89	0.87	83

Random Forest Classifier

```
In [21]: from sklearn.ensemble import RandomForestClassifier
```

```
In [22]: # Initialize the random forest classifier
rf_model = RandomForestClassifier(n_estimators=100, random_state=42)
```

```
In [23]: # Train the model on the training data
rf_model.fit(X_train, y_train)
```

```
Out[23]: ▼      RandomForestClassifier      ⓘ ?
RandomForestClassifier(random_state=42)
```

```
In [24]: # Make predictions on the test set
y_pred_rf = rf_model.predict(X_test)
```

```
In [33]: # Calculate the accuracy
accuracy = accuracy_score(y_test, y_pred_rf)
print(accuracy)
```

0.891566265060241

```
In [25]: # Confusion matrix
confusion_matrix(y_test, y_pred_rf)
```

```
Out[25]: array([[ 4,  9],
                [ 0, 70]])
```

```
In [26]: # Check accuracy of model
print(classification_report(y_test,y_pred_rf))
```

	precision	recall	f1-score	support
NO	1.00	0.31	0.47	13
YES	0.89	1.00	0.94	70
accuracy			0.89	83
macro avg	0.94	0.65	0.71	83
weighted avg	0.90	0.89	0.87	83

Support Vector Machine model

```
In [27]: from sklearn.svm import SVC
```

```
In [28]: # Initialize the Support Vector Machine model
svm_model = SVC(kernel='linear', random_state=42)
```

```
In [29]: # Train the model on the training data
svm_model.fit(X_train, y_train)
```

```
Out[29]: SVC
SVC(kernel='linear', random_state=42)
```

```
In [30]: # Make predictions on the test set
y_pred_svm = svm_model.predict(X_test)
```

```
In [34]: # Calculate the accuracy
accuracy = accuracy_score(y_test, y_pred_svm)
print(accuracy)
```

0.9397590361445783

```
In [31]: # Confusion matrix
confusion_matrix(y_test, y_pred_svm)
```

```
Out[31]: array([[ 8,  5],
                [ 0, 70]])
```

```
In [32]: # Check accuracy of model
print(classification_report(y_test,y_pred_svm))
```

	precision	recall	f1-score	support
NO	1.00	0.62	0.76	13
YES	0.93	1.00	0.97	70
accuracy			0.94	83
macro avg	0.97	0.81	0.86	83
weighted avg	0.94	0.94	0.93	83

## Model Comparison Report

Logistic regression model and Random Classifier have have same accuracy of 89%. Whereas Support Vector Machine model has accuracy of 94%. Therefore, Support Vector Machine model is best model for production.

This notebook was converted with [convert.ploomber.io](https://convert.ploomber.io)